



DRAFTLY

Grounded legal retrieval and a lawyer-controlled conveyancing matter workflow for Sri Lanka

GROUP 06 · UNIVERSITY OF MORATUWA · DCS&E

Himath Dhanapala · Lahiru Dilshan · Praveen De Silva

Supervisor: Dr. Nisansa de Silva



A conveyancing matter has two information bottlenecks

01

LEGAL INFORMATION RETRIEVAL

One scenario can require several provisions. A missing section leaves the answer incomplete.

02

MATTER DOCUMENTS

Key facts are buried in PDFs and scans. Lawyers must read and reconcile them.

MEASURED LEGAL SEARCH + SOURCE-LINKED DOCUMENT REVIEW



Practitioners shaped the corpus and the matter workflow



ANURA DHANARATNA

Lawyer · Notary · Law College lecturer

3 IN-PERSON CONSULTATIONS

Notarial curriculum

→ Topic map and statutory coverage

Sample matter documents

→ Document-reading trials

Practitioner feedback

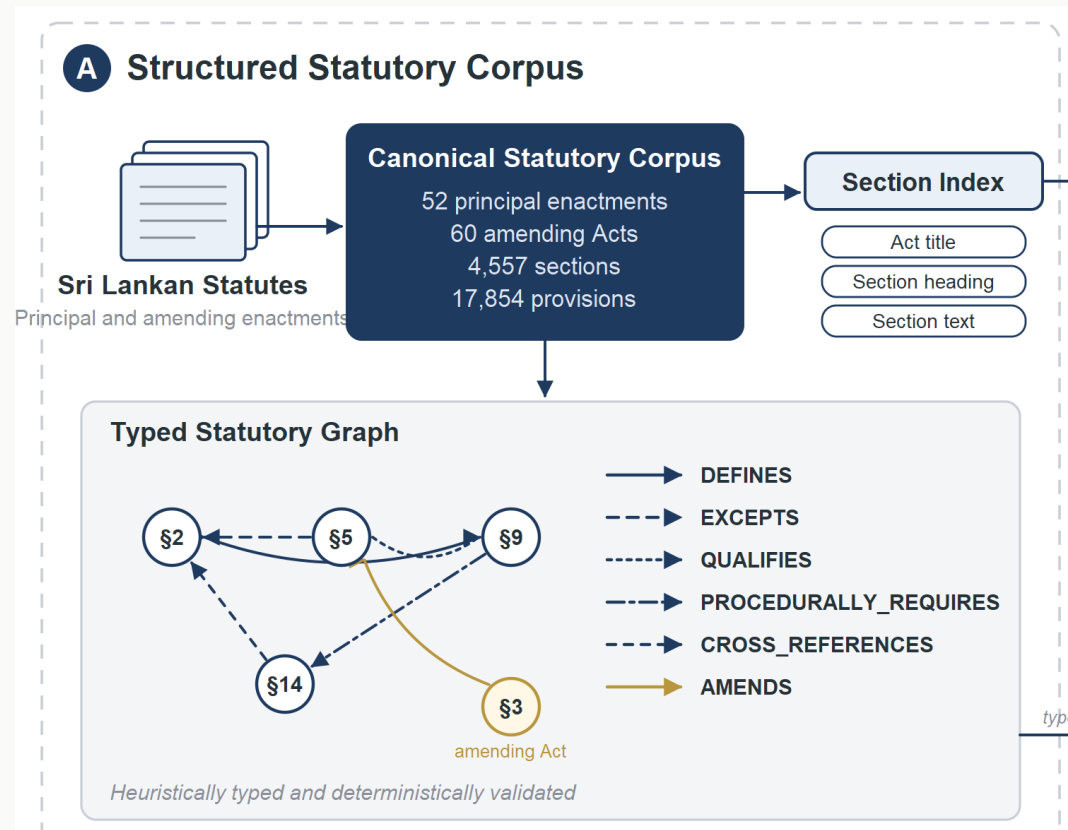
→ Evidence review and lawyer decision gates

ADDITIONAL STAKEHOLD
CONVERSATIONS





Law College course tables guided statutory corpus selection



52 + 60

principal enactments + amending Acts

4,557

sections · 17,854 provisions

6,346

typed statutory edges



Public Act source-link index · parsed text stays internal

INTERNAL PARSED RESEARCH CORPUS



The case-law collection adds judgments, with provisional rule extraction

3,703

reported NLR / SLR conveyancing judgments

5,474

official Supreme Court / Court of Appeal judgment documents

3,521 / 3,703

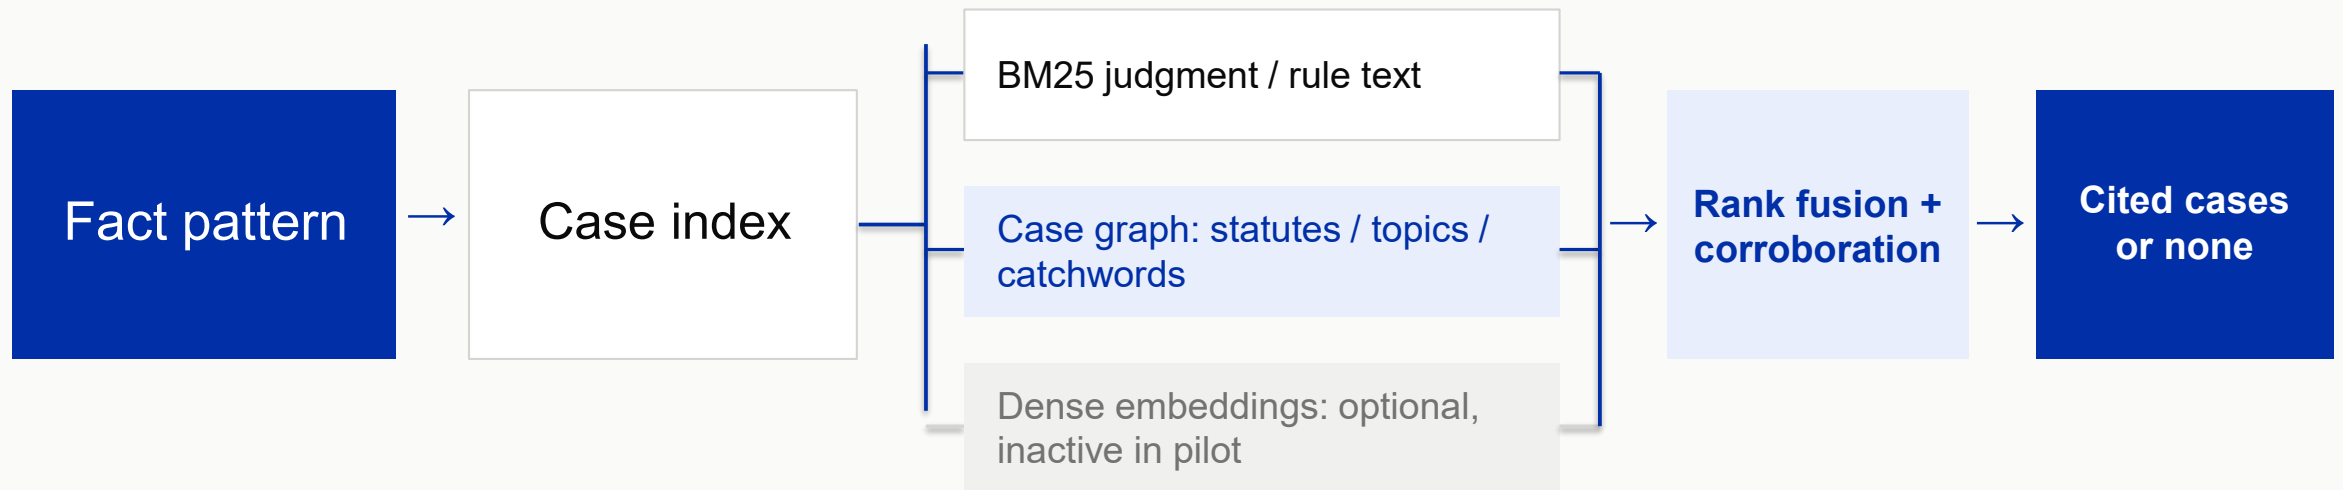
reported cases yielded verbatim headnote rules

PROVISIONAL RULES · LAWYER VALIDATION PENDING



A separate retriever finds analogous conveyancing cases

5,121 CONVEYANCING-FLAGGED CASES / 9,177 COLLECTED



20-query proxy pilot

14 / 20 on the initial corpus
snapshot

10 / 20 after rebaseline

LLM-judged relevance · no attorney labels



Law College scenarios became a provisional statutory QA benchmark

01 · QUESTION SELECTION

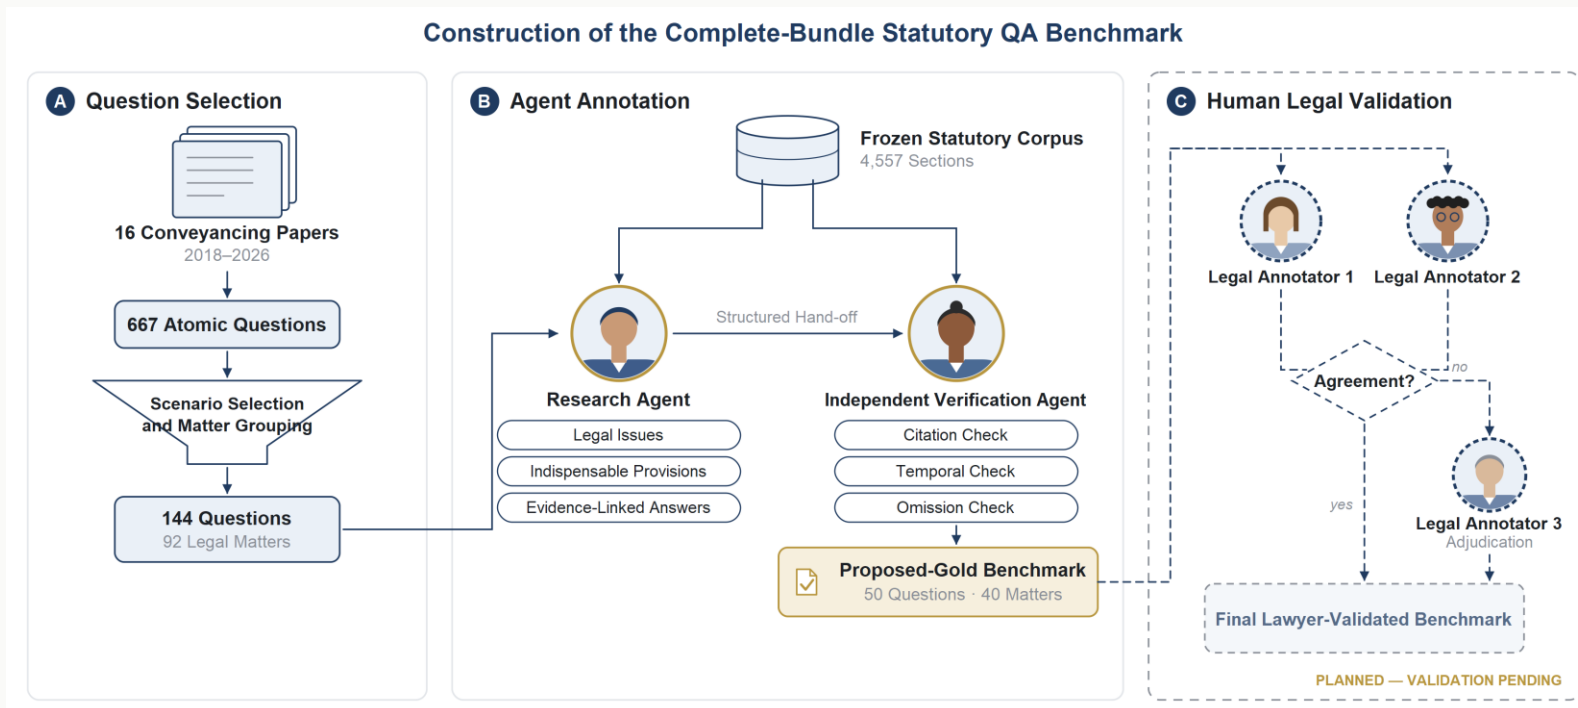
16 papers → 667 atomic questions

02 · PROVISION LABELS

50 scored questions / 40 matters

03 · HUMAN REVIEW

Three-attorney stage pending



40 TEST QUESTIONS / 32 MATTERS · PROVISIONAL LABELS

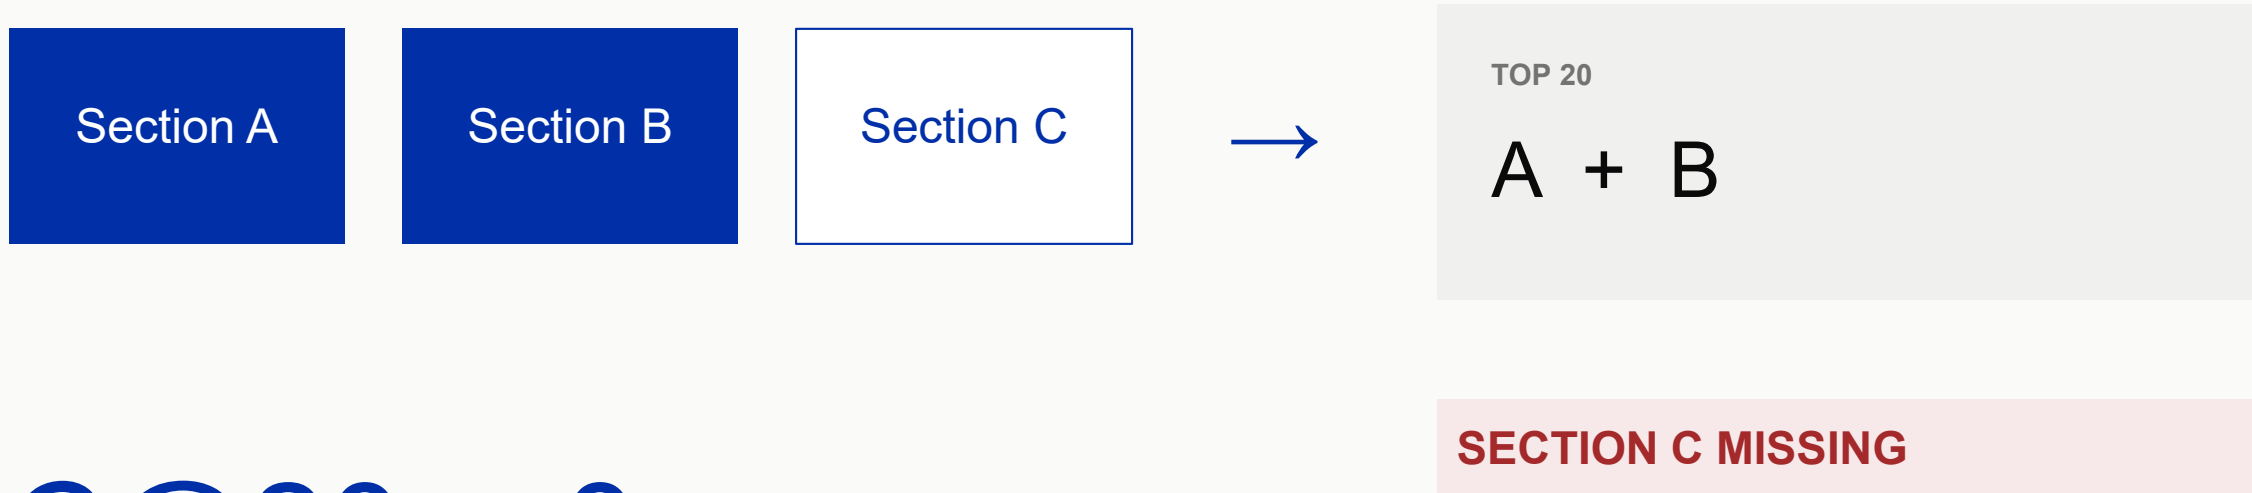


Provisional release



Retrieval succeeds only when it finds every proposed indispensable section

ILLUSTRATIVE QUESTION: AGENT-PROPOSED REFERENCE BUNDLE

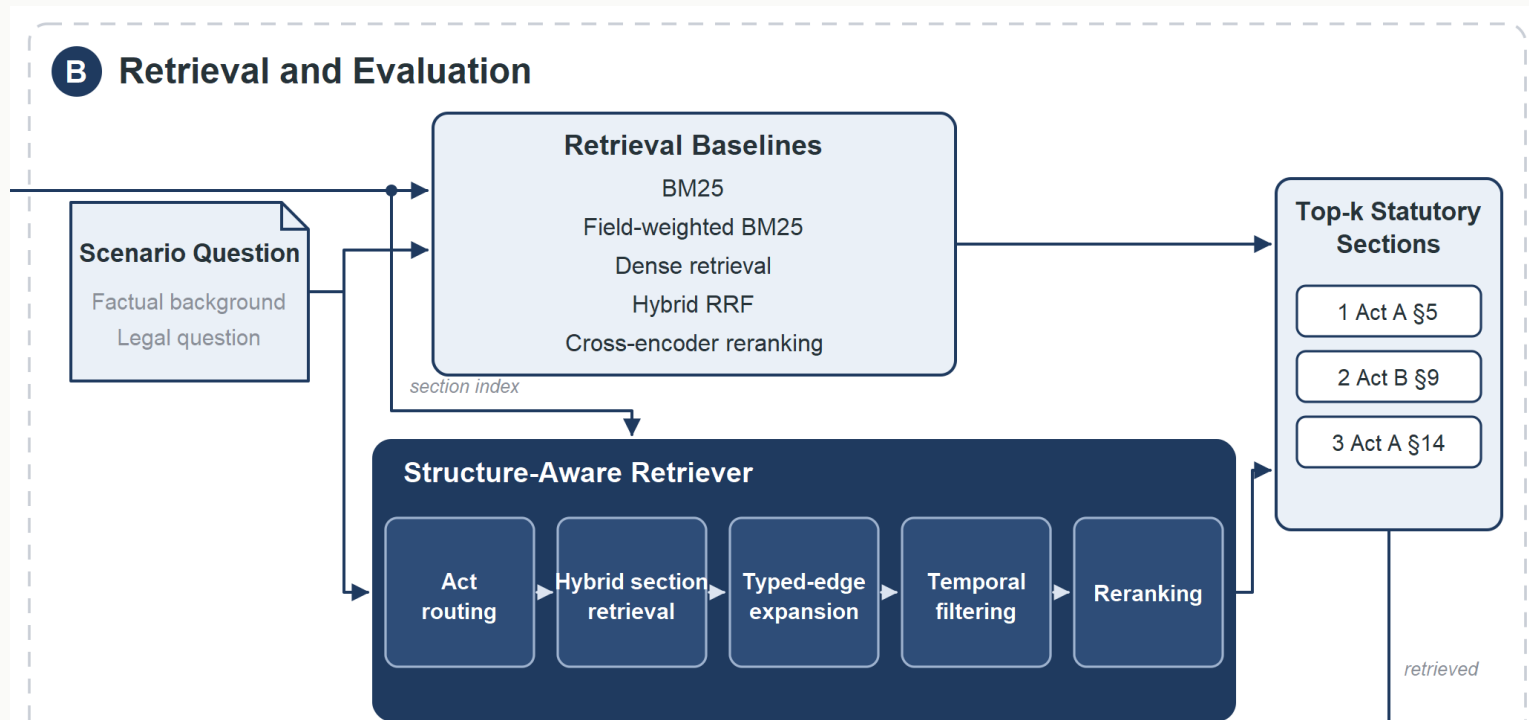


$$C@20 = 0$$

Full credit only when every proposed section appears in the first 20 results.



Seven offline retrieval systems used the same benchmark



LEXICAL

BM25 · field-weighted BM25

SEMANTIC / FUSION

Dense · hybrid · reranker

STRUCTURE

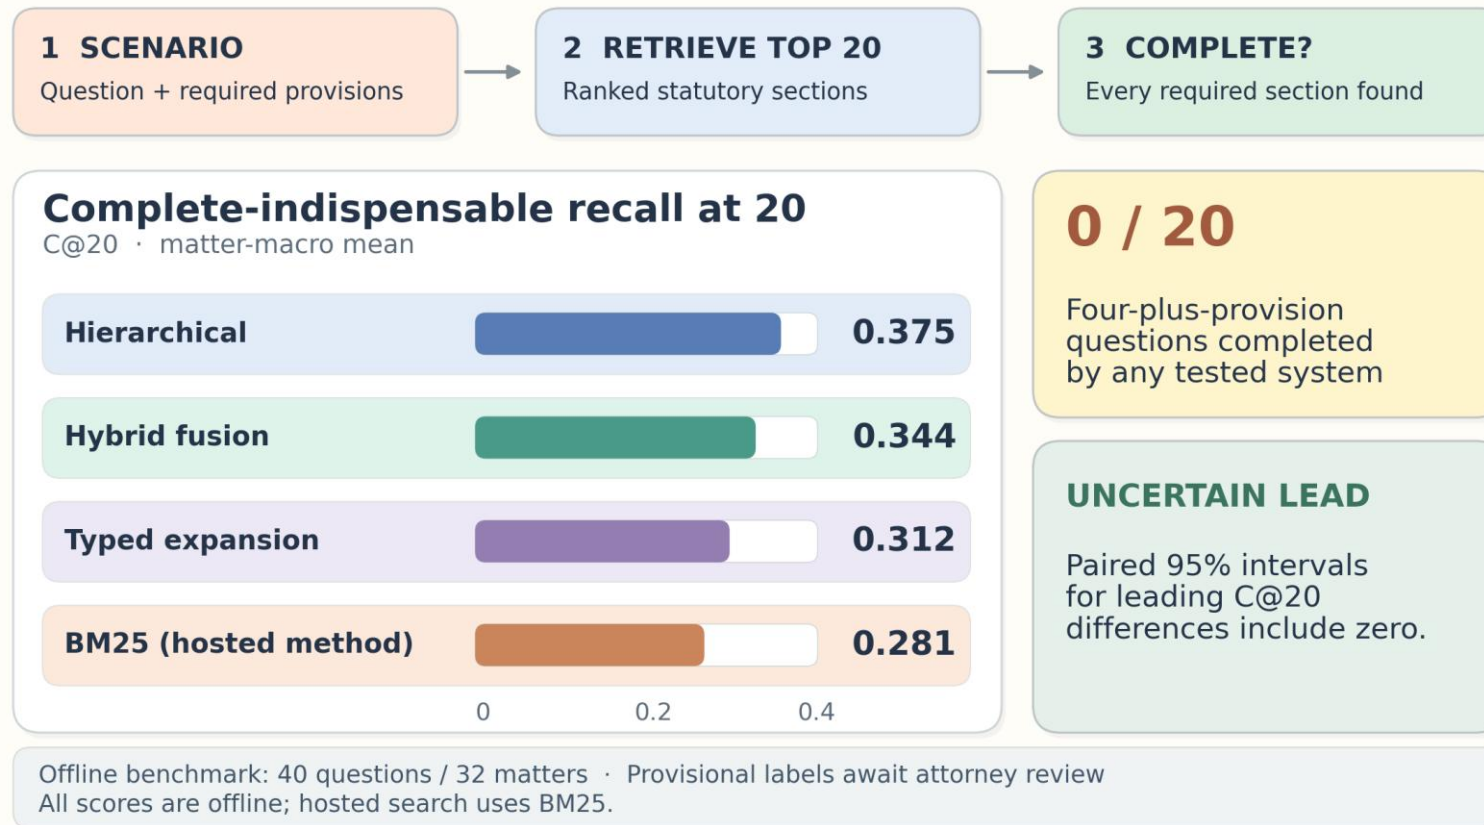
Hierarchical · typed expansion

OFFLINE RESEARCH · PROVISIONAL LABELS

Seven matched configurations · one frozen test split



Complete recall peaked at 0.375, while hard bundles remained unsolved



0 / 20

questions needing ≥4 provisions completed

0 / 20

Four-plus-provision
questions completed
by any tested system

UNCERTAIN LEAD

Paired 95% intervals
for leading C@20
differences include zero.

2 / 114

missed provisions one explicit edge away

No clear winner among the
leading systems

OFFLINE · PROVISIONAL LABELS · 40 TEST QUESTIONS / 32 MATTERS



The accepted paper reports the complete-bundle retrieval result

Does Statutory Structure Help? Retrieving Every Indispensable Provision for Sri Lankan Conveyancing Scenarios

Himath Dhanapala, Lahiru Dilshan, Praveen de Silva, Nisansa de Silva
Dept. of Computer Science & Engineering, University of Moratuwa, Sri Lanka
{himathd.23, lahirud.23, praveend.23, nisansadds}@cse.mrt.ac.lk

ACCEPTED FOR PRESENTATION

NeurIPS 2026 GlobalSouthAI workshop

01 Structured Sri Lankan statutory corpus

02 Provisional scenario QA benchmark

03 Seven-system retrieval evaluation

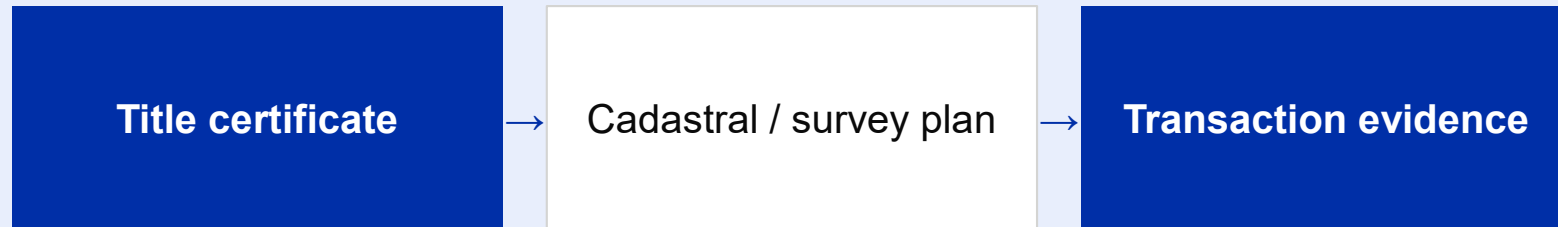
Typed-edge expansion did not demonstrate a complete-recall gain over hybrid retrieval.



Post-certificate RTA gives the first workflow a bounded scope

FIRST SCOPE · REGISTRATION OF TITLE ACT (RTA)

DEFERRED



Handwritten pattiru / folio
OCR

AT-form extraction

Historical title
reconstruction

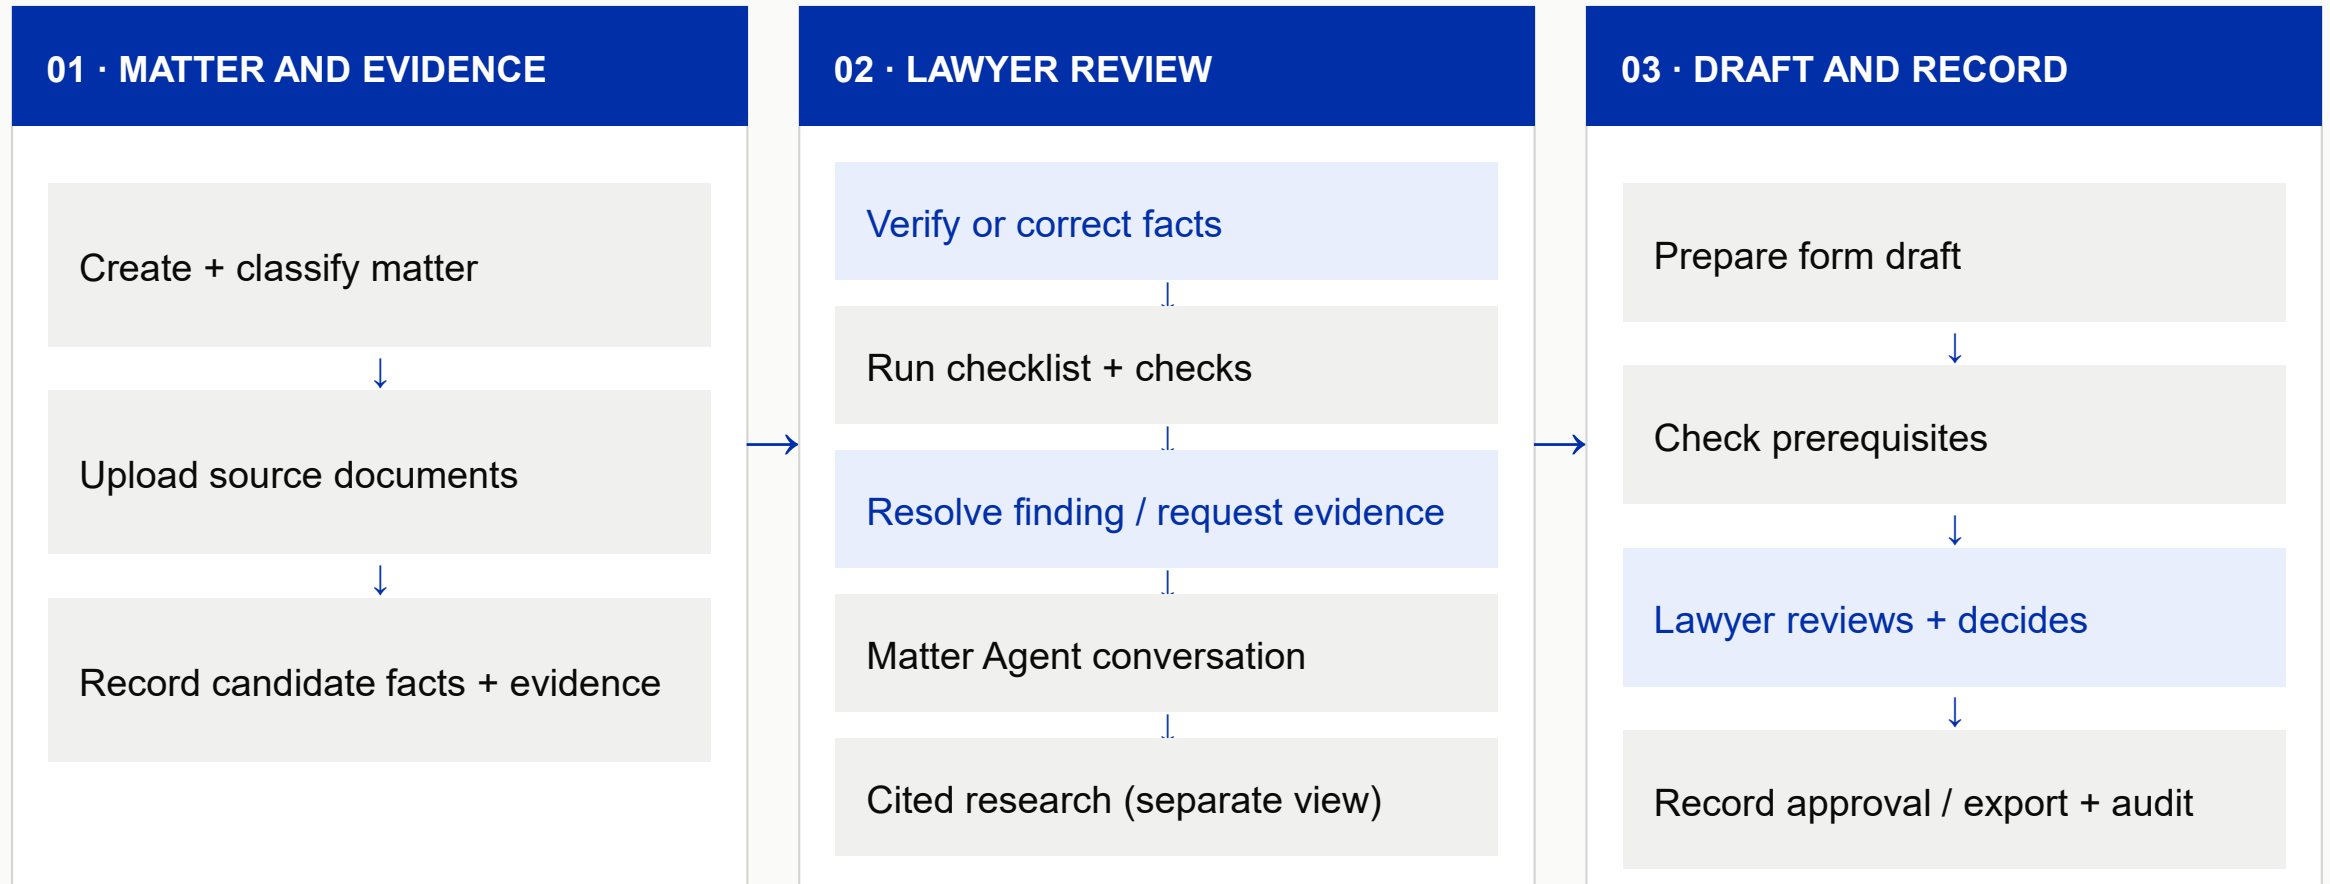
Initial title settlement

Defined requirements + prescribed forms

support reviewable checks



One matter connects intake, review, checks, and audit

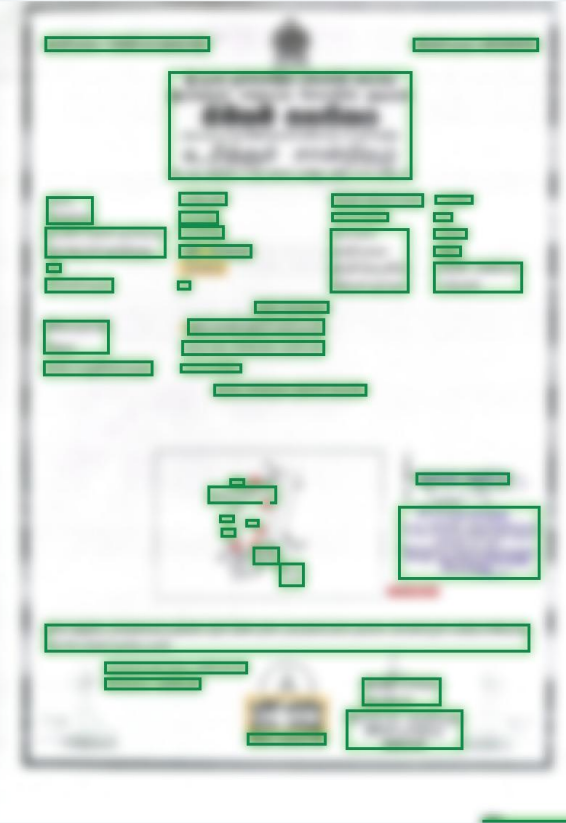


Findings return to review; a rejected draft returns to verified facts.



OCR detects text regions, but field accuracy remains unscored

REAL OCR REGION OVERLAY



Source text redacted; region boxes retained
Illustrates detection, not field accuracy

Real benchmark page · text redacted · OCR regions shown

4 matters · 38 documents

282 pages in the benchmark inventory

26 / 26 pages

pilot pages returned text

37 labelled fields

across 4 documents in 1 matter

FIELD-LEVEL ACCURACY NOT YET SCORED



The full processing design ends at lawyer review

01 · CAPTURE

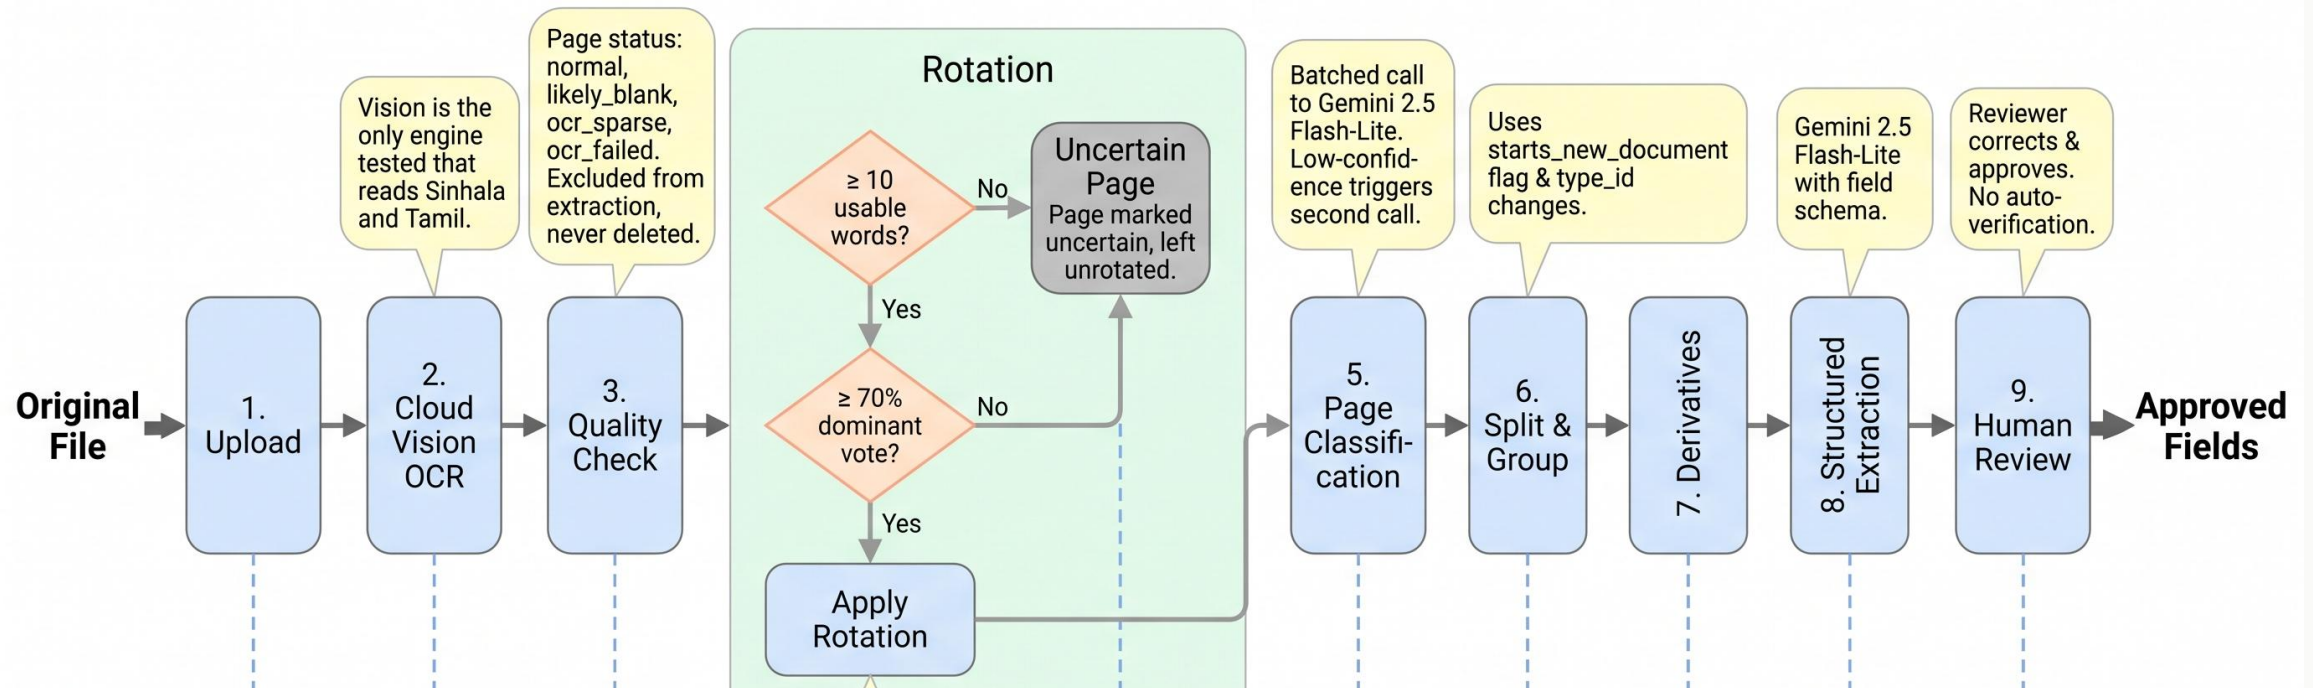
Upload · OCR · quality check

02 · ORGANIZE

Rotate · classify · group

03 · PROPOSE + VERIFY

Derivatives · fields · human review



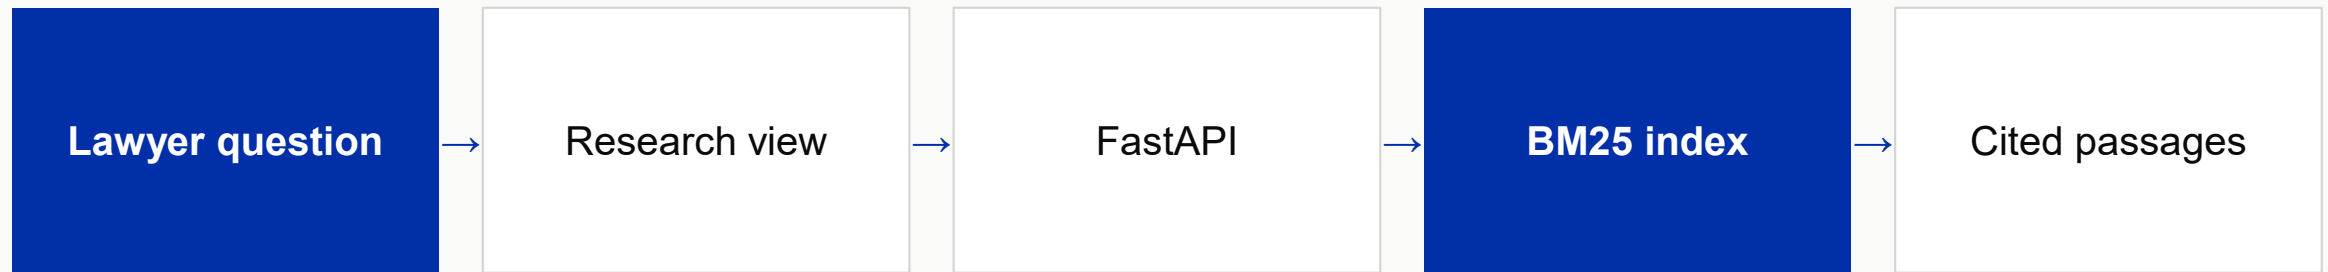
RESEARCH DESIGN · LATER EXTRACTION STAGES ARE NOT HOSTED

Hosted today: upload → processing record / stub candidates → lawyer review



The shipped research service returns cited statutory passages

HOSTED RESEARCH PATH · BM25



Grounded response
or insufficient authority

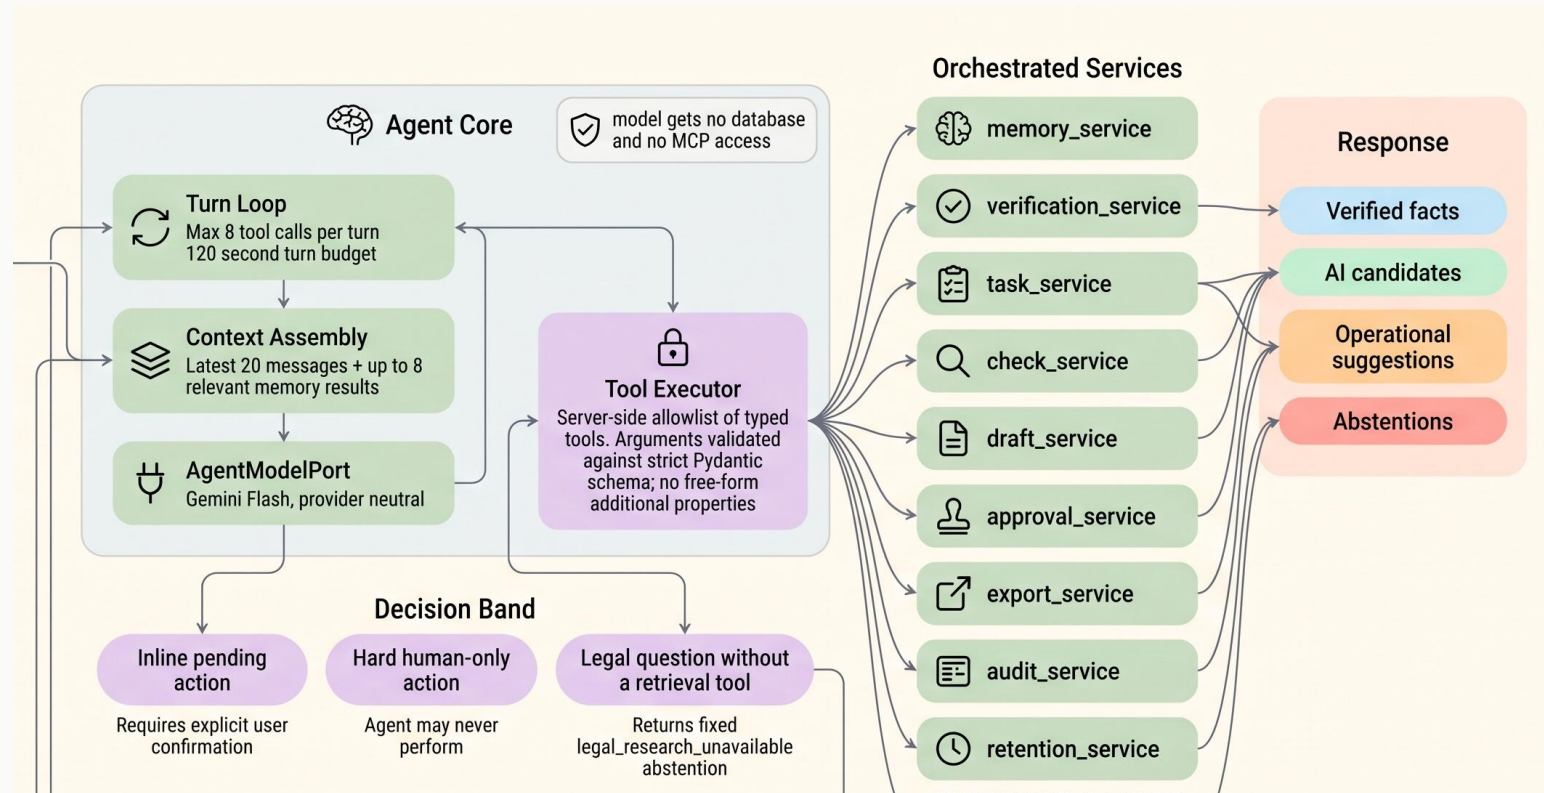
CURRENT BOUNDARY

Separate research workspace; no
direct Matter Agent search

Next version in progress: embeddings + query rewriting



The assistant can propose actions; the server and lawyer control them



Matter-scoped session

Allowlisted server tools

Action → explicit confirmation

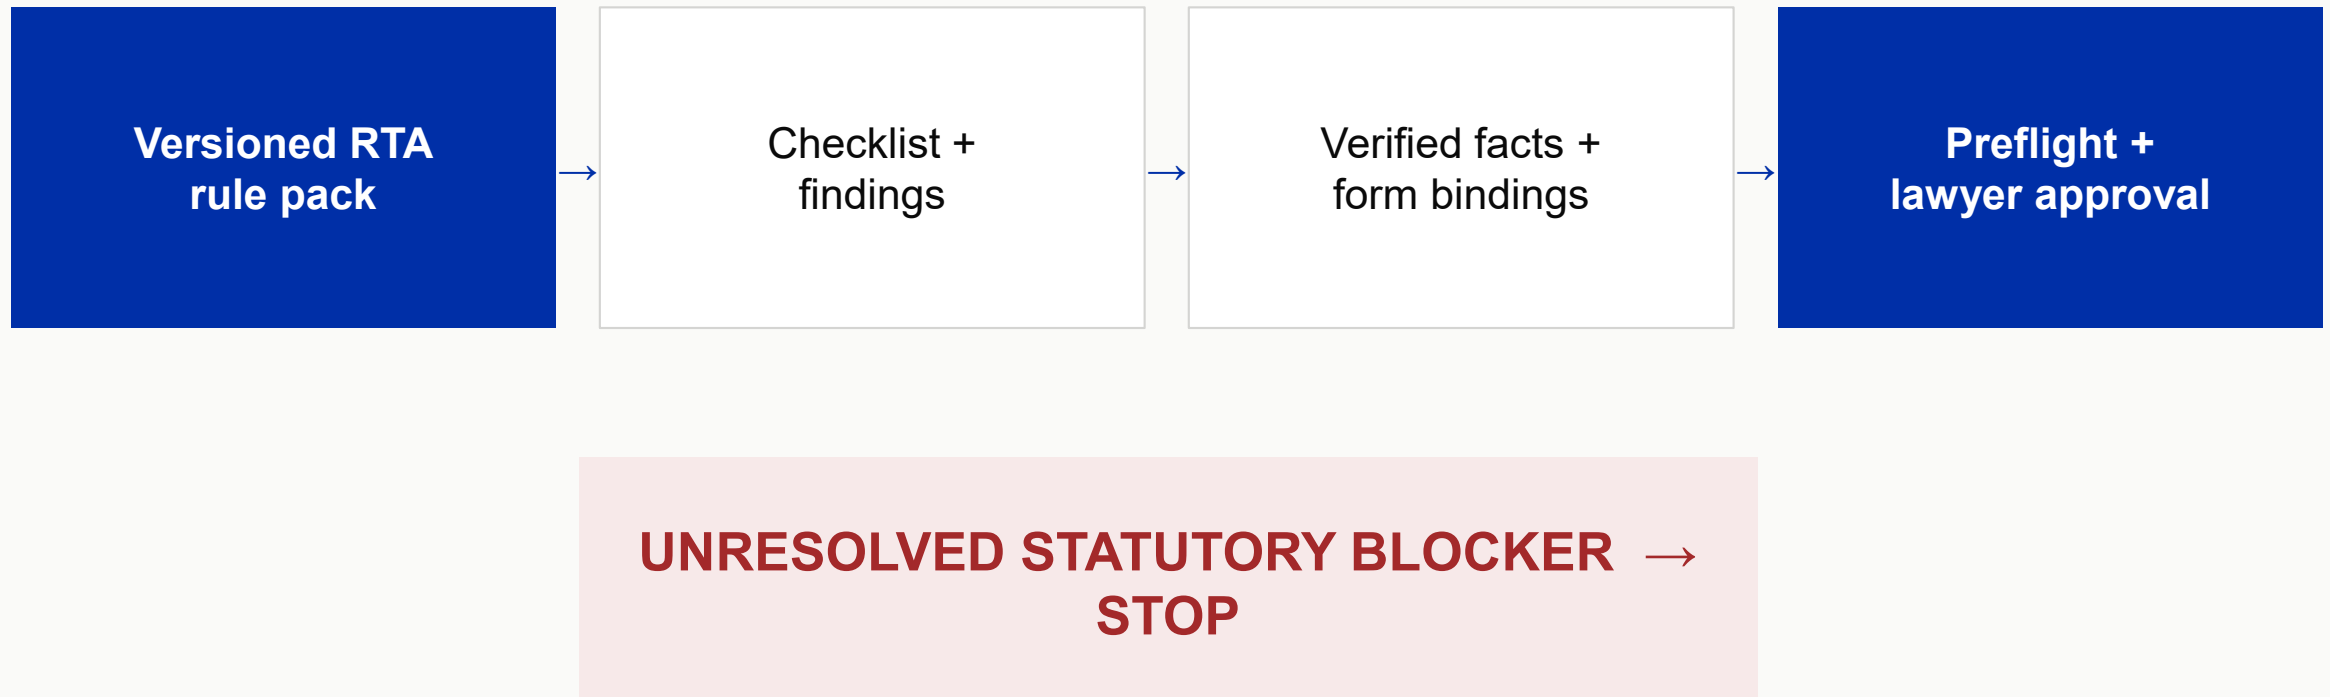
Legal question → abstention

Cited research → separate view

EARLIER DESIGN REFERENCE



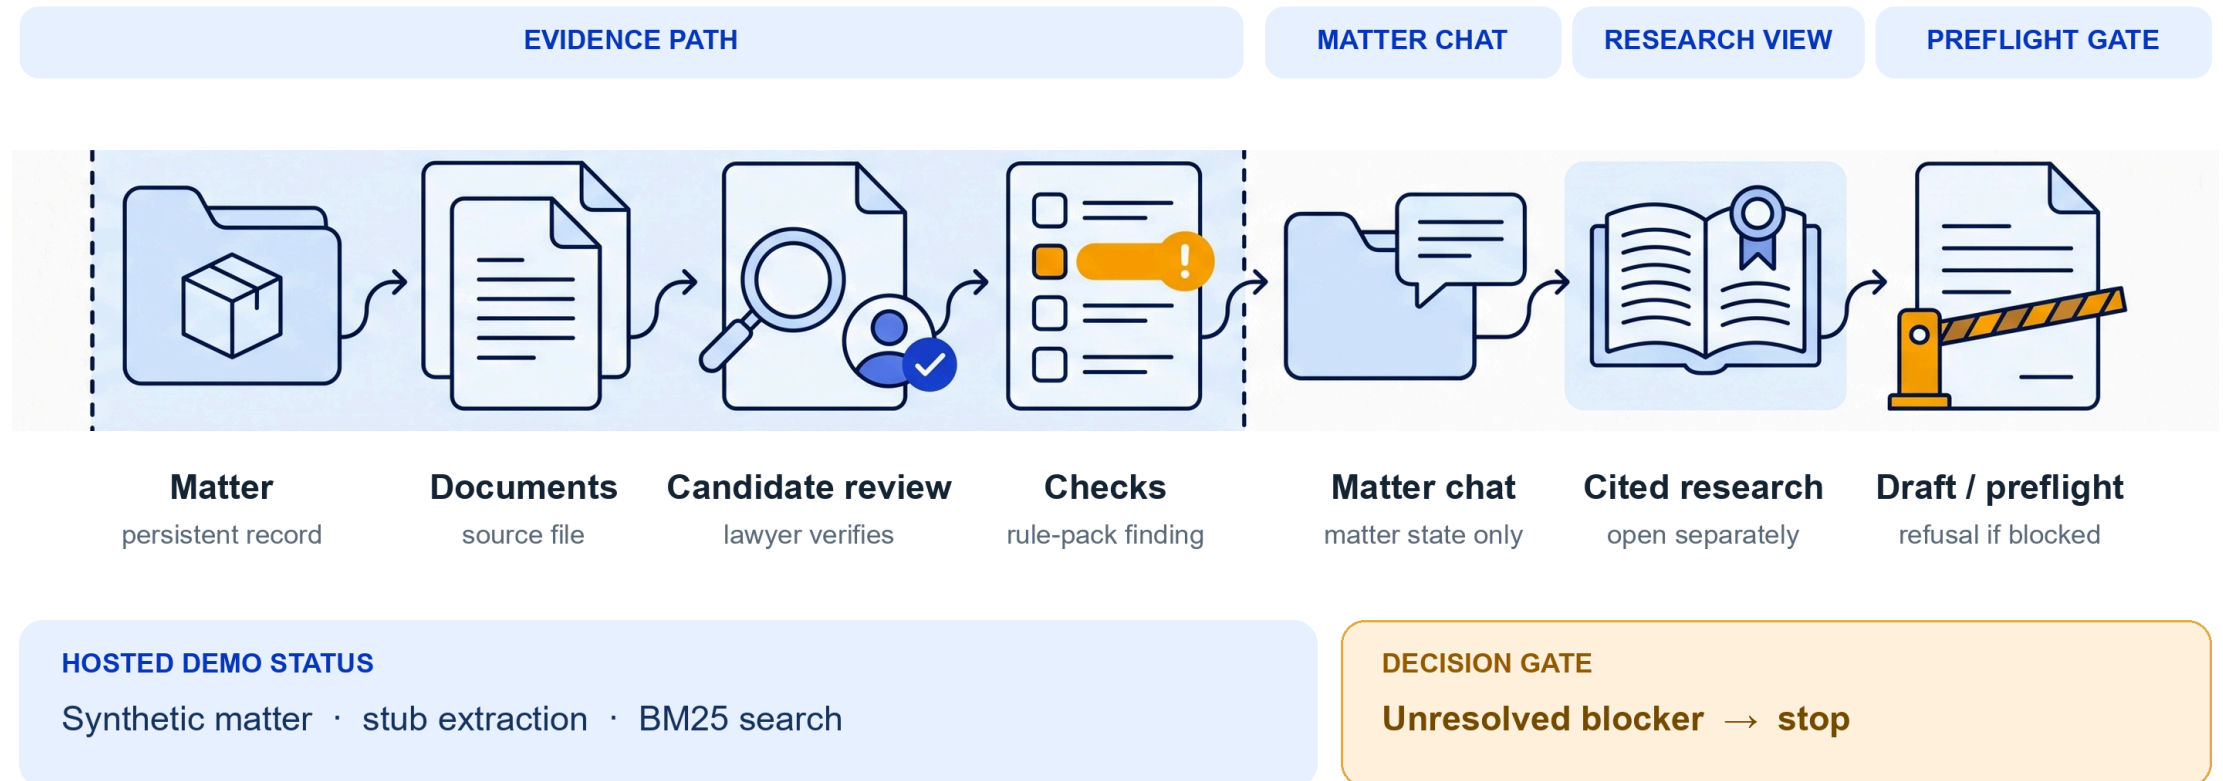
The rule pack can stop an unsafe drafting path



An unapproved draft cannot be exported as approved.



Inspect a synthetic RTA matter from evidence to refusal gate



Switch to the rehearsed live app or verified recording.



Five test tracks show what Draftly has measured

STATUTORY IR	C@20 0.375 · 40 test questions	Provisional labels
SIMILAR CASES	14/20 initial · 10/20 rebaseline	Proxy judged; dense inactive
CASE-LAW RULES	74/100 quote-grounded	58.1% proxy usable
DOCUMENT READING	26/26 pilot pages returned text	Field accuracy unscored
PLATFORM	3,906 backend + 348 frontend passed	3 browser refusals; one journey
WIDER BROWSER RUN: 4 PASSED · 9 FAILED · 2 SKIPPED		



Three contributors connected research and the product

HIMATH DHANAPALA

- Statutory corpus + benchmark lead
- Retrieval evaluation
- High Court case-law work
- Research-paper first draft

LAHIRU DILSHAN

- Case-law collection + retrieval
- OCR research
- Document-extraction research

PRAVEEN DE SILVA

- BM25 search work
- Templates + output flow
- Lawyer-facing review interface

SHARED: INTEGRATION, TESTING, AND FINAL PROJECT REVIEW



Measured retrieval.
Versioned evidence.
Lawyer-owned decisions.

NEXT VALIDATION

Attorney labels · document fields · form review · complete browser journeys

THANK YOU